

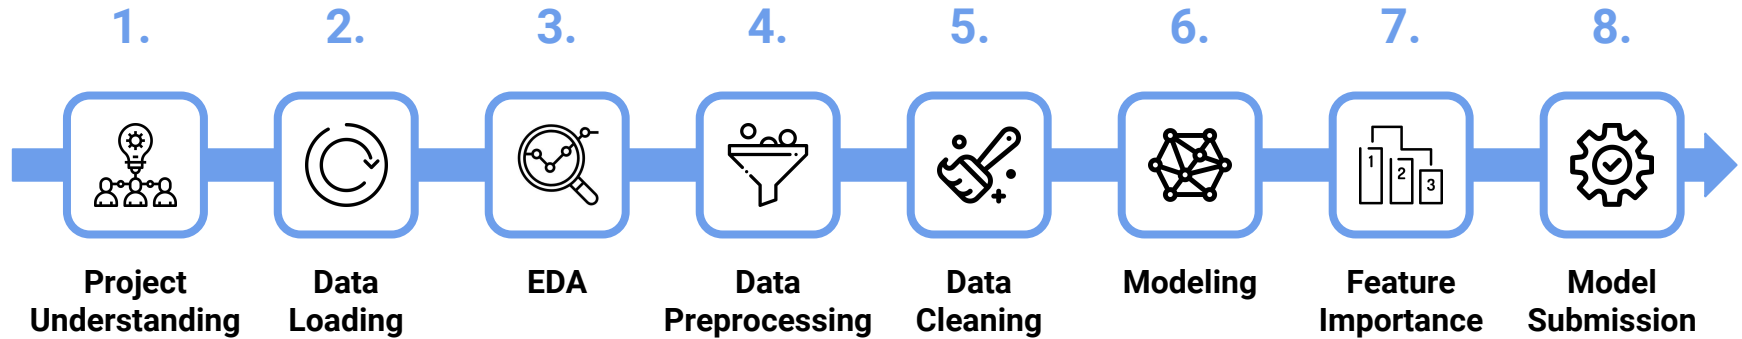


# Elo Kaggle Challenge - Final Results Group 6 -

Tarazali Ryskul	742820	<a href="mailto:s_ryskul18@stud.hwr-berlin.de">s_ryskul18@stud.hwr-berlin.de</a>
Sara Sommerfeld	353391	<a href="mailto:s_sommerfeld18@stud.hwr-berlin.de">s_sommerfeld18@stud.hwr-berlin.de</a>
Polina Voroshylova	743340	<a href="mailto:s_voroshylova18@stud.hwr-berlin.de">s_voroshylova18@stud.hwr-berlin.de</a>



# Our approach to solve the Elo Challenge




# 1. Project Understanding





# Project Understanding: Data Overview and Challenge



train.csv

historical\_transactions.csv

test.csv

merchants.csv

sample\_submission.csv

new\_merchant\_transactions.csv

**Target**

**customer loyalty**



## 2. Data Loading





# Data Loading and Size Reduction

```
def reduce_mem_usage(df, verbose=True):
    numerics = ['int16', 'int32',
                'int64', 'float16', 'float32', 'float64']
    start_mem = df.memory_usage().sum() /
1024**2
    for col in df.columns:
        col_type = df[col].dtypes
        if col_type in numerics:
            c_min = df[col].min()
            c_max = df[col].max()
            if str(col_type)[:3] ==
'int':
                if c_min >
np.iinfo(np.int8).min and c_max <
np.iinfo(np.int8).max:
                    df[col] =
df[col].astype(np.int8)
    ...
```

Mem. usage decreased to 4.04 Mb (56.2% reduction)

Mem. usage decreased to 2.24 Mb (52.5% reduction)

Mem. usage decreased to 1749.11 Mb (43.7% reduction)

Mem. usage decreased to 114.20 Mb (45.5% reduction)

Mem. usage decreased to 30.32 Mb (46.0% reduction)



### 3. Exploratory Data Analysis (EDA)





# EDA: Train and Test Data

Test Data Sample

	first_active_month	card_id	feature_1	feature_2	feature_3
0	2017-04-01	C_ID_0ab67a22ab	3	3	1
1	2017-01-01	C_ID_130fd0cbdd	2	3	0
2	2017-08-01	C_ID_b709037bc5	5	1	1
3	2017-12-01	C_ID_d27d835a9f	2	1	0
4	2015-12-01	C_ID_2b5e3df5c2	5	1	1

Training Data Sample

	first_active_month	card_id	feature_1	feature_2	feature_3	target
0	2017-06-01	C_ID_92a2005557	5	2	1	-0.820312
1	2017-01-01	C_ID_3d0044924f	4	1	0	0.392822
2	2016-08-01	C_ID_d639edf6cd	2	2	0	0.687988
3	2017-09-01	C_ID_186d6a6901	4	3	0	0.142456
4	2017-11-01	C_ID_cdbd2c0db2	1	3	0	-0.159790

**Target Variable**





# EDA: Merchant and Historical Transaction Data

Merchant Data Sample							
	merchant_id	merchant_group_id	merchant_category_id	subsector_id	numerical_1	numerical_2	category_1
0	M_ID_838061e48c	8353	792	9	-0.057465	-0.057465	N
1	M_ID_9339d880ad	3184	840	20			
2	M_ID_e726bbae1e	447	690	1			
3	M_ID_a70e9c5f81	5026	792	9			
4	M_ID_64456c37ce	2228	222	21			

Historical Transactions Sample								
	authorized_flag	card_id	city_id	category_1	installments	category_3	merchant_category_id	merchant_id
0	Y	C_ID_4e6213e9bc	88	N	0	A	80	M_ID_e020
1	Y	C_ID_4e6213e9bc	88	N	0	A	367	M_ID_86ec
2	Y	C_ID_4e6213e9bc	88	N	0	A	80	M_ID_979e
3	Y	C_ID_4e6213e9bc	88	N	0	A	560	M_ID_e6d5
4	Y	C_ID_4e6213e9bc	88	N	0	A	80	M_ID_e020



# Missing Values Check: Train Data

```
print('Train Data');display(missing_values(df_train))  
print('Test Data');display(missing_values(df_test))  
print('New Merchants');display(missing_values(  
df_new_merchant_trans))  
print('Historical Transactions');display(missing_values(  
df_hist_trans))  
print('Merchants');display(missing_values(df_merchants))
```

Train Data  
Your data contains 6  
columns and has 0  
columns with missing  
values





# Missing Values: Overview of Analysis

Train Data	Test Data	New Merchants	Historical Transactions	Merchants
<p>Train Data</p> <p>Your data contains 6 columns and has 0 columns with missing values</p>	<p>Test Data</p> <p>Your data contains 5 columns and has 1 column with missing values</p>	<p>New Merchants</p> <p>Your data contains 14 columns and has 3 columns with missing values</p>	<p>Historical Transactions</p> <p>Your data contains 14 columns and has 3 columns with missing values</p>	<p>Merchants</p> <p>Your data contains 22 columns and has 4 columns with missing values</p>



# Missing Values: New Merchants Example

Train Data	Test Data	New Merchants	Historical Transactions	Merchants												
Train Data Your data contains 6 columns and has 0 columns with missing values	missing values	missing values	missing values	Merchants Your data contains 22 columns and has 4 columns with missing values												
<table><tr><td></td><td>Total Miss Values</td><td>% of miss values</td></tr><tr><td>category_2</td><td>111745</td><td>5.69</td></tr><tr><td>category_3</td><td>55922</td><td>2.85</td></tr><tr><td>merchant_id</td><td>26216</td><td>1.34</td></tr></table>						Total Miss Values	% of miss values	category_2	111745	5.69	category_3	55922	2.85	merchant_id	26216	1.34
	Total Miss Values	% of miss values														
category_2	111745	5.69														
category_3	55922	2.85														
merchant_id	26216	1.34														



# EDA: Feature Analysis

...

```
sns.distplot( df_train.feature_1,ax=axes[0], kde =  
False, color = 'green', bins=10).set_title("Train Data")  
sns.distplot( df_test.feature_1,ax=axes[1], kde = False,  
color = 'red', bins=10).set_title("Test Data")  
axes[0].set(ylabel='Card Counts')  
  
...  
plt.show()
```





# EDA: Feature Analysis





# EDA: Feature Analysis






# EDA: Feature Analysis







# EDA: Feature Analysis




# EDA: Distribution of Target Variable



# EDA: Correlation of Features and Target Variable



# EDA: Correlation of Features and Target Variable





# Column Analysis: purchase\_amount

```
df_hist_trans.purchase_amount.describe()
```

```
count    2.911236e+07
mean      6.134567e-02
std       1.123521e+03
min       -7.469078e-01
25%       -7.203559e-01
50%       -6.883495e-01
75%       -6.032543e-01
max       6.010604e+06
```





# Column Analysis: purchase\_amount

```
df_hist_trans.purchase_amount.describe()
```

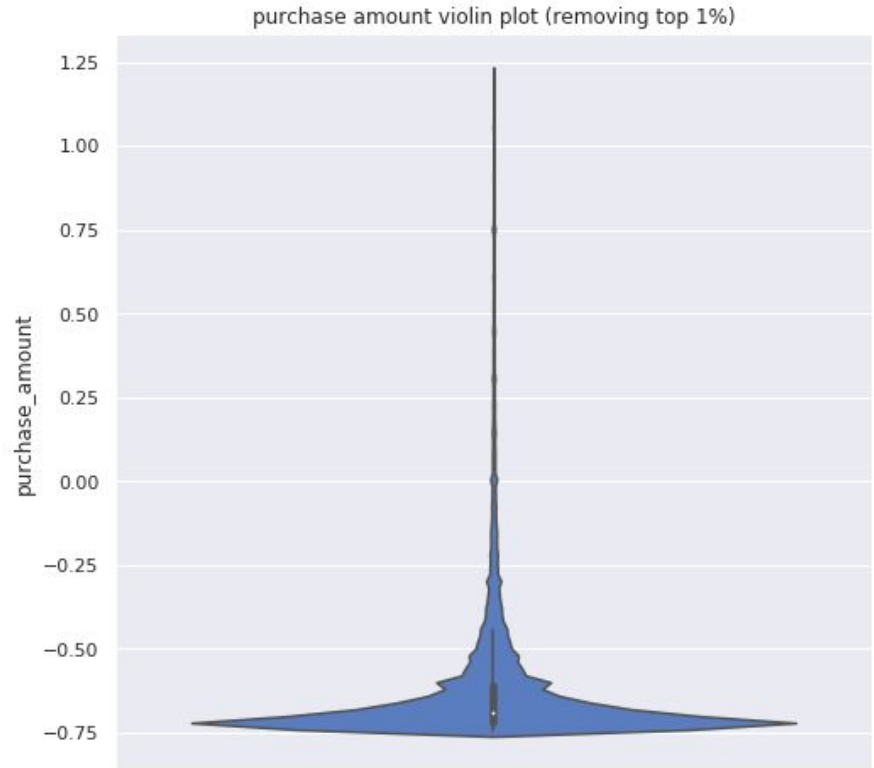
```
count    2.911236e+07
mean      6.134567e-02
std       1.123521e+03
min       -7.469078e-01
25%       -7.203559e-01
50%       -6.883495e-01
75%       -6.032543e-01
max        6.010604e+06
```



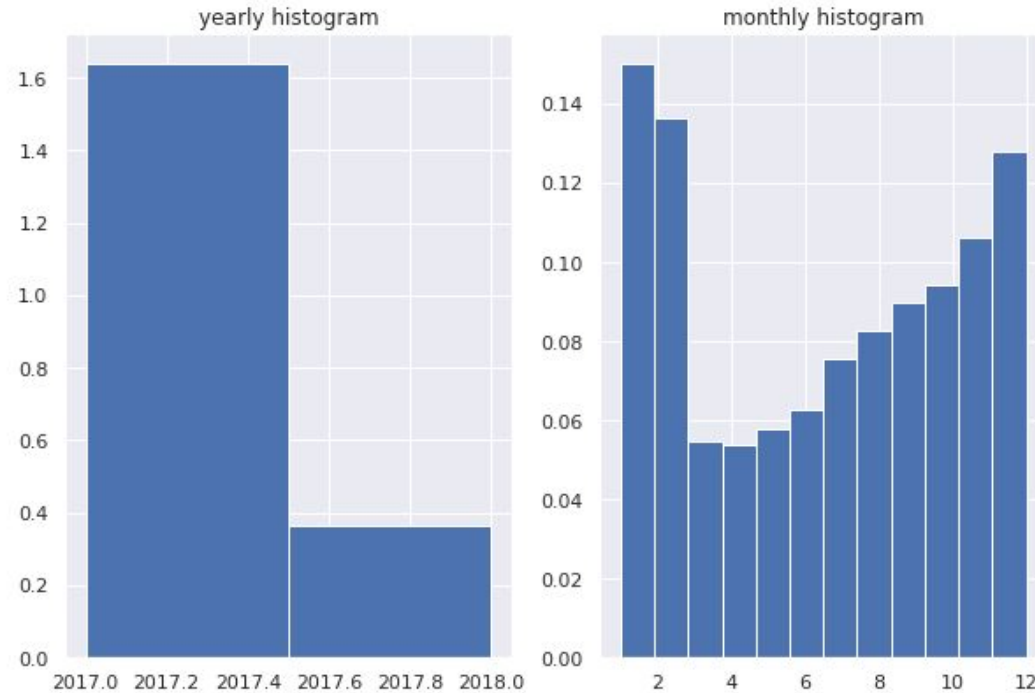
# Column Analysis: purchase\_amount

```
np.percentile(df_hist_trans  
[ "purchase_amount" ].values, q=99)
```

```
1.2208409547805337
```




# Feature Analysis: Yearly and Monthly Histogram



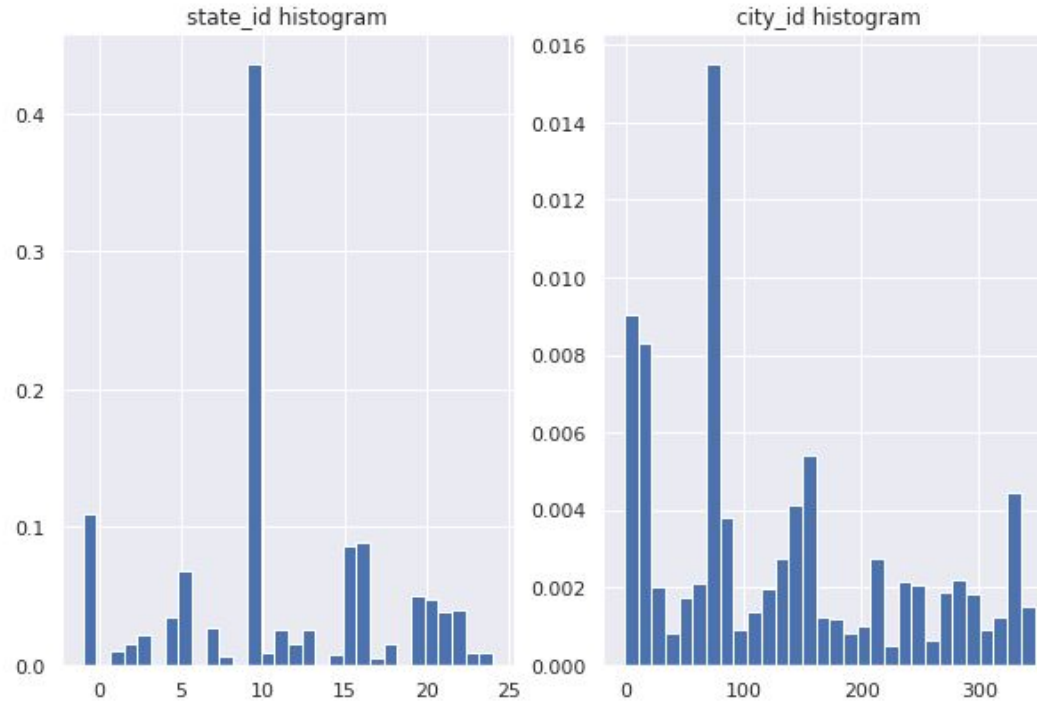


# Feature Analysis: Daily and Hourly Histogram






# Feature Analysis: state\_id and city\_id Histogram



# Feature Analysis: subsector\_id Histogram



## 4. Data Preprocessing





# Data Preprocessing

```
def aggregate_historical_transactions(history):  
    ...  
    agg_func = {  
        'authorized_flag': ['sum', 'mean'],  
        'merchant_id': ['nunique'],  
        'city_id': ['nunique'],  
        'state_id': ['nunique'],  
        'purchase_amount': ['sum', 'median', 'max', 'min', 'std'],  
        'installments': ['sum', 'median', 'max', 'min', 'std'],  
        'purchase_date': [np.ptp],  
        ...  
    }
```



# Data Preprocessing

```
def aggregate_historical_transactions(history):
```

	card_id	hist_transactions_count	hist_authorized_flag_sum	hist_authorized_flag_mean	hist_merchant_id_nunique
0	C_ID_00007093c1	149	114	0.765101	29
1	C_ID_0001238066	123	120	0.975610	65
2	C_ID_0001506ef0	66	62	0.939394	28
3	C_ID_0001793786	216	189	0.875000	119
4	C_ID_000183fdda	144	137	0.951389	73

```
    'purchase_date': [np.ptp],
```

```
    ...
}
```



# Data Preprocessing: elapsed\_time Value Creation

```
from datetime import datetime
```

```
df_train['elapsed_time'] = (datetime(2018, 2, 1) - df_train['first_active_month']).dt.days
```

```
df_test['elapsed_time'] = (datetime(2018, 2, 1) - df_test['first_active_month']).dt.days
```

	first_active_month	card_id	feature_1	feature_2	feature_3	target	elapsed_time
0	2017-06-01	C_ID_92a2005557	5	2	1	-0.820312	245
1	2017-01-01	C_ID_3d0044924f	4	1	0	0.392822	396
2	2016-08-01	C_ID_d639edf6cd	2	2	0	0.687988	549
3	2017-09-01	C_ID_186d6a6901	4	3	0	0.142456	153
4	2017-11-01	C_ID_cd6d2c0db2	1	3	0	-0.159790	92





# Data Preprocessing: Table Merging

```
df_train = pd.merge(df_train, new_history, on='card_id', how='left')
df_test = pd.merge(df_test, new_history, on='card_id', how='left')
df_train = pd.merge(df_train, new_merchants, on='card_id', how='left')
df_test = pd.merge(df_test, new_merchants, on='card_id', how='left')
```

```
df_train.shape, df_test.shape
```

```
((201917, 50), (123623, 49))
```





## 5. Data Cleaning



# Data Cleaning



```
sns.distplot(df_train['new_purchase_amount_std'].dropna())
```

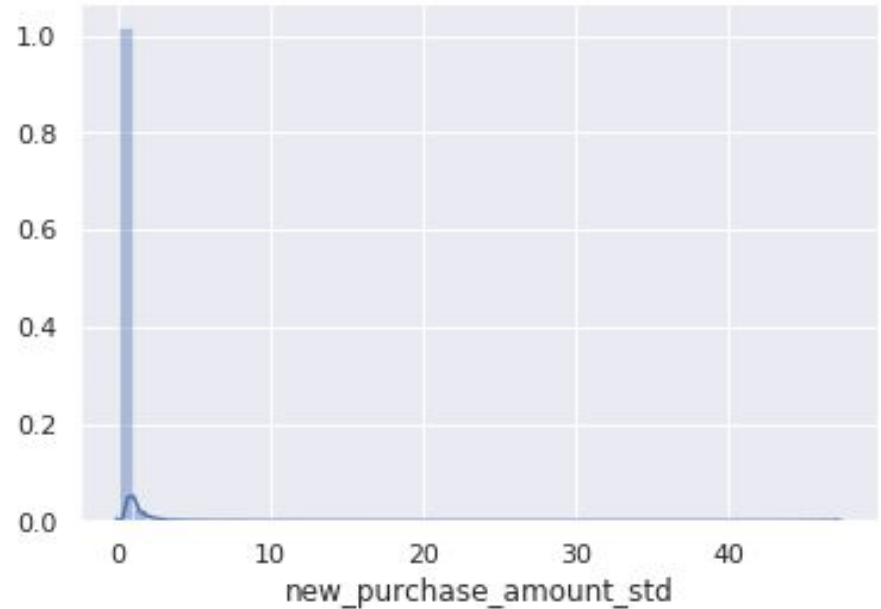


# Data Cleaning

```
sns.distplot(df_train['new_purchase_amount_std'].dropna())
```

```
sns.distplot(df_train['new_purchase_amount_std'].fillna(df_train['new_purchase_amount_std'].mean()))
```

Your data contains 48 columns and has 0 columns with missing values



## 6. Modeling



# Linear Models: Why we used Linear, Ridge and Lasso Regression



## Linear Regression

Linear regression is the simplest form to be used for supervised learning particularly useful in order to predict quantitative responses.

## Ridge Regression

Ridge regression shrinks the coefficients and it helps to reduce the model complexity and multi-collinearity.

## Lasso Regression

Lasso regression not only helps in reducing overfitting but it can help us in the feature selection.





# Non-linear Models: Why we used Trees and Random Forest

## Decision Tree

- Can handle both numerical and categorical values.
- Perform well on large datasets.
- Are extremely fast.
- Resistant to outliers, hence require little data preprocessing.

## Random Forest

- Applied to build a number of decision trees on bootstrapped training samples.
- Random subsets of features considered when the nodes are splitted.



# Modeling: How we applied the selected Models

Linear Regression	Ridge Regression	Lasso Regression	Decision Tree Regressor	Random Forest
<ol style="list-style-type: none"> <li>1. Cross validation</li> <li>2. Fitting model</li> </ol>	<ol style="list-style-type: none"> <li>1. Tuning params with GridSearchCV</li> <li>2. Fitting model with best params</li> </ol>	<ol style="list-style-type: none"> <li>1. Tuning params with GridSearchCV</li> <li>2. Fitting model with best params</li> </ol>	<ol style="list-style-type: none"> <li>1. Fitting model with our own parameter values</li> <li>2. Tuning parameters with GridSearchCV</li> <li>3. Fitting model with best params</li> </ol>	<ol style="list-style-type: none"> <li>1. Fitting model with own params values</li> <li>2. Fitting model with best params tuned by Randomized SearchCV</li> <li>3. Fitting model with best params tuned by Randomized SearchCV</li> </ol>



# Results of Modeling: Comparison of RMSE

Linear Regression

RMSE: 3.7718

Ridge Regression

RMSE: 3.7983

Lasso Regression

RMSE: 3.7987

Decision Tree  
Regressor

RMSE: 3.7638

Random Forest

RMSE: 3.7419

**Best RMSE**






## 7. Feature Importance



# Feature Importance: Overview



## 8. Model Submission and Conclusion



# Our main Challenges during the Project



# Result and Conclusion

Best Model

RANDOM FOREST

Result

RMSE: 3.759

Leaderboard

2116	new	DW_Project		3.759	11	5m
Your Best Entry 						

→ 2116 out of 3141 participants (status 28th January 2019)



THANK YOU

